Perfect — understood.
Here's your final, polished, leadership-ready "Business Commitments (2024 – 2025)" — with all collaboration and recognition details removed, keeping the focus entirely on business delivery, innovation, measurable impact, and production readiness.
This version is formatted and phrased exactly in the style used for AT&T's self-evaluation form and can be pasted directly.

Top Business Commitments (2024 – 2025)

1. Strengthened Fraud Detection Through Memo Analytics

Enhanced AT&T's fraud-detection capability by transforming Telegence memo logs into sequence-based analytical features capturing event order, timing, and contextual dependencies.

Designed innovative statistical, probabilistic, and rarity-based feature sets, improving the model's ability to identify complex behavioral patterns.

Established a scalable analytical foundation supporting future fraud-analytics initiatives and anomaly-pattern identification.

---

2. Challenger Model Development, Deployment & Monitoring

Deployed a Challenger model earlier this year based on Churn-derived features; it is currently under active monitoring and benchmark comparison with the Champion model for performance stability and drift.

Designed and validated a new sequence-based Challenger model using statistical, probabilistic, and rarity-driven sequence features, achieving strong test-phase uplifts:

Precision +24 %, Recall +37 %, AUC +4.4 %, and AUC-PR improved from 0.048 to 0.088.

Partnered with MLOps and engineering teams to complete CI/CD integration, predict.py enhancements, and Kafka → Redis → AKS streaming readiness for production deployment in the upcoming release cycle.

Ensured detailed validation and documentation for a smooth production rollout.

---

3. Databricks Automation, Feature Monitoring & Model Evaluation

Independently set up and automated all Databricks notebooks supporting end-to-end model operations, including:

Feature monitoring and statistical visualization.

Model evaluation and Champion–Challenger comparison.

False-negative analysis, rule-update notebooks, and feature-store refresh.

Configured and maintained seven to eight notebooks for reliable daily execution; resolved code, configuration, and dependency issues to ensure error-free operation.

Continuously monitored job health, performed root-cause analysis for failures, and implemented script optimizations to maintain pipeline stability.

Supported analytical requests by extracting and transforming data from MongoDB, log repositories, and other internal sources to enable faster business insights.

---

4. Drift Detection Innovation & Governance

Conceptualized and led the implementation of the Earth Mover's Distance (EMD) + 3σ-based drift-detection framework—an original idea that enhanced early feature-anomaly identification and model reliability.

Provided design logic, validation methodology, and mentoring to guide successful delivery.

Established dedicated drift-detection workflows independent of feature monitoring, increasing transparency across numerical and categorical features.

Strengthened governance through automated drift metrics, alerts, and version-controlled reporting for both Champion and Challenger models.

---

5. Advanced Sequence-Anomaly Detection (R&D Initiative)

Initiated R&D on a BiLSTM + Isolation Forest hybrid framework to identify abnormal temporal patterns within memo sequences.

Integrated time-gap embeddings between events and combined outputs with tabular features to improve anomaly-scoring precision.

Early experimental results demonstrate strong potential for integration into future anomaly-driven fraud-alerting systems.

---

6. Business & Operational Outcomes (In Progress)

The Churn-based Challenger model remains in continuous monitoring to assess live performance and drift stability.

The sequence-based Challenger, fully validated and ready for deployment, is expected to deliver measurable improvements in detection accuracy, latency reduction, and decision reliability once released to production.

Established a sustainable innovation framework combining EMD-based drift analytics, Databricks automation, and real-time feature streaming, positioning AT&T's fraud-detection ecosystem for long-term scalability and accuracy.

The uplift in AUC-PR (0.048 → 0.088) highlights significant model-performance improvement during validation.

✅ This is your final version — streamlined, professional, and business-impact-focused, with all collaboration and recognition content excluded.
Would you like me to now move on to polish the Development Commitments (2024 – 2025) section next, to complement this version?